

Market & Problem Landscape

PRODUCT CONTEXT

Claude offers powerful capabilities - Agents, Skills, Projects (Workflows), and Prompts. Each one solves a different class of problem.

But the user's real challenge is not capability access. It's a routing problem:

"I know what I want to do. I just don't know which Claude to use for it."

PROBLEM SPACE

Despite multiple capability tiers, choosing the right Claude mode still feels surprisingly stressful for most Pro users. People are not struggling to find capabilities.

They are struggling to feel confident the chosen mode will actually work - that it fits their task, won't waste tokens, and won't trigger a configuration spiral.

Today's tools treat capability choice as a search problem, not a decision-confidence problem.



MARKET TRENDS

- Global AI Assistant market: ~USD 41B by 2030
- AI Agent market: ~USD 28B by 2030, CAGR 43%
- Key driver: Demand shifts from chat to autonomous execution
- Gap: No player has built a routing-intelligence layer yet



CURRENT TOOLS AND WHY THEY FAIL

Tool	What They Do	Why They Fail
Cursor	Chat + Agent + Composer modes	No intent routing; users self-select and fail
ChatGPT GPT Store	3M+ custom GPTs to choose from	No routing intelligence; discovery collapses
Perplexity	Search + Pro Search + Spaces	Search-only routing; no agent or task tier
GitHub Copilot	Chat + Agent + Workspace modes	Power users adapt; mainstream defaults to chat
v0 by Vercel	UI generation only	Single-mode tool; no decision layer needed
Replit Agent	Autonomous coding only	Removes flexibility; one-size-fits-all execution
Devin / Manus	Autonomous agent only	No mode choice; users can't dial down autonomy
LangChain	Code-first agent framework	No help with "should I even use this?"

Primary & Secondary Research



PRIMARY RESEARCH

Survey : 25 Claude Pro/Team users · Behavioral questions about real past tasks · No hypotheticals

88%

failed first try
or gave up

64%

started prompting
without Agent setup


68%

retreated to basic
prompt after failure

12%

felt fully confident
in their choice (5/5)

Where exactly does the decision break down?



>88% of difficulty is in decision-making, not capability access

Source - [User Survey](#)



SECONDARY RESEARCH

AI capability adoption & decision psychology:

- **Multi-mode AI triggers feature-avoidance under uncertainty** : Users default to the simplest mode (NN/g)
- **Paradox of Choice in SaaS** : 3+ equally-prominent options take 40% longer to decide; 25% higher abandonment (HBR / Schwartz)
- **AI adoption readiness gap** : 53% of workers don't know how to get value from AI; 70% would use it more if they knew when to use it (Salesforce 2024)
- **Confidence > Accuracy as adoption driver**
Decision confidence - not output accuracy - drives AI usage frequency (MIT Sloan Management Review)
- **Causal chain for AI adoption**:
Reasoning → Certainty → Trust → Use intention (NUS / Microsoft Research, 2025)

Source - [In-depth research](#) →

User Segment & Pain Points



BEHAVIOR-BASED SEGMENTATION

- Unaware & Unbothered → *No pain*
- Aware & Unbothered → *No motivation*
- Unaware & Trying → *Hard to activate*
- **Aware & Frustrated** → **Sharp pain + Active attempts + Repeated failure (TARGET)**



WHY THIS SEGMENT WINS

- They already feel the problem (40% rate task-fit confusion as the #1 hardest decision factor)
- They've tried everything - Agents (44%), looked up docs (20%), asked the community (4%)
- They actively seek change but keep retreating to simple prompts (68% after first failure)

Arjun, 29 yrs
Senior Engineer | Bangalore

Profile

High-performing engineer, technically capable. Uses Claude daily for code review, debugging, research.

Pattern

Knows what task he wants automated → has clear goals → hits the configuration wall → spirals into trial-and-error → re-prompts manually 8 times.

Tried

Agent setup (gave up at config), Project setup (too rigid), custom prompts (works but exhausting), asked Discord.

Emotion

Quiet frustration and self-doubt rather than confidence. Blames himself: "I know Claude can do this - why can't I figure this out?"

What he wants

A pre-filled scaffold and a clear signal that the chosen capability fits his specific task - not more documentation.



HYPOTHESES TO TEST

Users don't lack intent - they lack a routing layer. They have clear tasks but can't translate them into a confident capability choice, and the **capability they try first shapes their outcome more than the task itself**. They feel this friction but misread it as personal failure, not a product gap. The fix isn't more options or better docs - it's showing users why a specific capability fits their specific task, in their own words. **Reasoning beats choice.**



RESEARCH INSIGHTS (N=25)

- 40% said "not knowing if Agent fits my task" was the hardest decision factor
- 48% delay or abandon advanced capability attempts despite having time
- Only 12% feel fully confident (5/5) when selecting a capability

Insights from interviews

"I know exactly what I want Claude to do. I just freeze when I have to pick the right mode." - Arjun

Detailed Problem Breakdown

WHAT IS THE TRUE PROBLEM?

Capability choice in Claude is meant to be empowering, yet most Pro/Team users get stuck in a painful gap: they know what they want to automate, but cannot confidently turn that knowledge into a capability choice that fits. The result is constant overthinking, last-minute manual workarounds, and settling for simple prompts.

HOW DO WE KNOW THIS IS REAL?

Survey users tell the same story repeatedly. They try an Agent, hit the configuration wall, ask the community, and still feel unsure even after switching modes. 88% fail or give up on first attempt. 68% retreat to basic prompts. Only 12% feel fully confident at the moment of choice. The pattern is universal.

WHO IS FACING THIS?

The Capable Explorers - Pro and Team users who genuinely want to leverage Claude's full capability suite. They are technically reflective and attentive in daily work, yet become anxious and indecisive the moment they face the Agent vs Project vs Prompt decision. They want to ship, but the pressure of getting it right leaves them defaulting to the safest path.

WHAT VALUE IS GENERATED BY SOLVING IT?

For the user - they move from frustrated guesswork to calm confidence. They finally use Claude to its full potential.

For Anthropic - Navigator becomes the trusted decision layer for capability choice. Users who adopt Agents successfully churn 3-5× less than those who don't. Direct retention + monetization lever.

WHY SHOULD WE SOLVE THIS PROBLEM NOW?

The capability suite is already built. The bottleneck is decision confidence - and every week this persists, high-intent users settle for simple prompting permanently. AI routing intelligence is a learnable moat: outcome data from every user decision compounds. Early data advantage is durable and defensible - but only if Anthropic builds the routing layer before competitors do.

Ideation and Prioritization / Job-to-be-done (JTBD)

SOLUTION 1: BETTER DOCS & TOOLTIPS

User hovers over capability options → richer explanations and glossary appear.

Fatal Flaw: Solves comprehension, not decision. Users already know what Agents do. They can't tell if it fits their specific task.

More information = more paralysis.

SOLUTION 2: COMMUNITY RECOMMENDATION LAYER

User shares task with community → upvoted patterns guide capability choice.

Fatal Flaw: Maps to real behaviour (some users do ask the community) but adds dependency, latency, and noise.

Doesn't build the user's own confidence.

SOLUTION 3: NAVIGATOR - INTENT ROUTING ENGINE (CHOSEN FOR MVP)

Start with the user's actual task → analyze intent → score all capabilities → show fit reasoning + tradeoffs → guide configuration.

Key Strength: The only solution that starts where users actually are (clear task + self-doubt) and directly solves the decision-confidence wall.



COMPARATIVE DECISION MATRIX

Criterion	Sol 1	Sol 2	Sol 3 (MVP)
Addresses task-fit gap (40%)	✗	Partial	✓
Works without external dependency	✓	✗	✓
Operates at the decision moment	✗	✗	✓
Builds the user's own confidence	✗	✗	✓
Scales across task types	✓	Partial	✓
RICE Score	28	35	160

Detailed Description for [Ideation](#)



JOB-TO-BE-DONE (JTBD)

When I need... to use Claude for a complex task, I want to...
translate my task into a confident capability choice, so I feel...
confident the chosen mode will succeed first try.




WHY NAVIGATOR IS THE MVP

It is the only solution that directly tests the core hypothesis from primary research: if we show users why a capability fits their task - with reasoning rooted in their own words - confidence increases and they will act.

Wireframing & Prototype (1)

1




HOME / GET STARTED

Clean, calm entry point. Tagline: "Describe your task - Navigator finds the right Claude mode."

Three value pillars set expectations:

- Reasoning, not lists
- Fit Score (the core promise)
- No config wall

2




DESCRIBE YOUR TASK

User types the task in plain language. No mode selectors, no dropdowns.

Navigator's complexity detector parses the input as the user types - fires only on multi-step tasks.

3




INTENT ANALYSIS

Navigator extracts 5 intent dimensions transparently:

- Task type
- Complexity score
- Tool needs
- Repetition flag
- Branching logic

First signal that Navigator understands the task, not just keywords.

4



PROJECT CONTEXT (OPTIONAL)

Lightweight personalisation. Optional. User pastes a Project ID, repo name, or context.

Used only to make later reasoning feel grounded - never required.


TOKEN BUDGET

Sets realistic expectations early. User picks a session budget tier.


Prevents over-routing later and anchors the tradeoff scoring.

Wireframing & Prototype (2)


5




6



7



8



EXISTING APPROACH? (OPTIONAL)

The differentiator. User can paste their existing prompt or workflow they're already considering - or skip.

This is where Navigator proves it starts with the user's own thinking, not a blank slate.

TELL US ABOUT THE TASK

Guided 5-question intent dimensions panel.

Questions extract task structure the user already knows but hasn't articulated.

Every answer is reflected verbatim in the fit reasoning later - this is the moat.

HERE'S WHAT WE RECOMMEND

Core output screen.

Shows the user's existing approach (if entered) with a fit score + reasoning.

Then presents 3 ranked Capability Cards with "Why it fits" - quotes user's exact words.

User swipes to compare. No config yet.

YOUR PICK

User accepts a capability.

Shows a clean summary with a risk flag and "How to make it stronger" note.

Includes the predicted token cost, output quality, and latency.

CONFIGURATION SCAFFOLD (PAID TIER)

Paid-tier unlock.

Pre-filled Agent / Project / Skill configuration with sensible defaults.

Saved templates, route history, and outcome learning - turns the recommendation into action without forcing a config wall.

Customer Journey Mapping

USER JOURNEY : NAVIGATOR ROUTING

1. The Confusion

Action: Opens Claude, types complex task. Reverts to prompting.

"I know Claude can do this - why can't I figure this out?"

2. The Routing

Action: Navigator intercepts. Extracts 5 intent dimensions.

"It's actually reading what I want to do - not keywords."

3. The Decision

Action: Accepts recommendation (92% fit) with confidence.

"This explains why. The reasoning tracks."

4. The Confidence

Action: Agent executes successfully. Saves as template.

"I chose this - and I know why - and it worked."

[LINKS](#)

SYSTEM DESIGN LAYERS

L1: Data Ingestion

Task desc, Project context, Token budget



L2: Signal Processing

Intent vector, Complexity classifier



L3: AI Engine (The Brain)

LLM API call, Capability scorer



L4: Output Deliverable

Fit score, Reasoning, Tradeoff bars



L5: Storage & Learning

Browser-local, Feedback loop

Business Case & Monetization Plan



THE METRICS HIERARCHY

We measure confidence formed, not sessions completed.

North Star: Task-Fit Confidence Uplift · Target: >60% of users report confidence ≥4/5 post-recommendation

Why it matters: If users complete sessions but don't feel more certain, the product has failed regardless of any engagement metric.

Level	Metric	Formula	Target	Why It Matters
L0	Session Completion Rate	Sessions reaching capability pick ÷ sessions started	>60%	Is the flow compelling?
L0	Routing Acceptance Rate	Sessions accepting top recommendation ÷ recommendations shown	>65%	Are routings resonant?
★ NS	Confidence Uplift Rate	% reporting ≥4/5 confidence post-session	>60%	Direct measure of core job
L1	3-Month Return Rate	Returning users ÷ total users	>30%	Value retained across tasks
L1	Token Efficiency Gain	Tokens vs baseline (no Navigator)	-40%	Paid-tier pull - does Navigator save cost?
Guard	Generic Routing Rate	% reporting recommendation felt arbitrary	<15%	Trust signal - too generic = product failure



MONETIZATION PLAN

Free Tier (Core Validation)

- Unlimited Navigator sessions
- Confidence scoring + capability recommendation
- Goal: prove the thesis that reasoning builds confidence

Paid Tier - Pro/Team

- Pre-filled config scaffold + saved templates
- Token spend estimates + outcome learning
- Route history across sessions and Projects
- Goal: turn validated confidence into effortless action

Future B2B (Phase 2)

- Anthropic Console integration for enterprise teams
- Routing analytics + capability adoption dashboards
- Willingness-to-pay pre-validated: 76% of survey users said explicit reasoning would change their confidence in using Agents

GTM & Risks/Mitigation



DISTRIBUTION & GO-TO-MARKET

Principle First

*Make the Capable Explorer feel seen before they've even used Navigator.
The product should feel specific, reasoned, non-generic.*

Top 3 Channels & Activation

1 In-Product Contextual Trigger (Primary, zero cost)

When input complexity threshold fires (≥ 3 intent signals), Navigator surfaces inline: "Before you send this, here's what will work best."
Reaches every user without marketing spend. Highest leverage.

2 Power User Community Seeding (Launch week)

Claude Discord and r/ClaudeAI have Capable Explorers by definition. Seed Navigator access to top contributors. They test hard, give signal-rich feedback, and amplify organically. Cost: low. Signal quality: high.

3 Email Activation for Dormant Pro Users

Users who pay for Pro but haven't used Agents are leaving value on the table. One specific data-backed email per user: "You use Claude for [X]. Here's how Navigator cuts that time in half with Agents."



RISKS AND MITIGATION

Risk	Why It Matters	Mitigation	Kill Switch
Recommendation feels generic	Generic = trust destroyed = product failure	Anchor every recommendation to user's exact words. No card ships without a quoted task fragment.	If >20% report generic → rebuild prompt architecture
Over-routing to Agents	Unexpected token bills create user backlash	Recommend Agent only when ≥ 3 intent signals fire. Show cost estimate before execution.	If >25% override → recalibrate thresholds
Users distrust AI advising AI	Meta-trust problem blocks adoption	Show auditable reasoning: "4 steps, web search, conditional logic - Agent"	If trust score <3/5 in 30 days → redesign reasoning display
Episodic use → low LTV	If users return only for hard tasks, product won't compound	Routing memory layer drives cross-task return. Saved templates keep Navigator top-of-mind.	If 3-month return <20% → build memory layer earlier
Cold start - task is vague	When intent signals are ambiguous, Navigator can't route confidently	Surface clarifying question rather than guessing. Default to safest capability tier.	Monitor ambiguity rate as leading indicator